

SEMINAR PROGRAM • C05151

Two seminars, sixteen topics

Every group presents one core topic in Seminar 1 (Weeks 3–6) and one advanced topic in Seminar 2 (Weeks 8–11): a 30-minute talk plus a mandatory 15-minute Q&A/defense. Seminar 1 and Seminar 2 are each 25% of the final grade.

30' + 15' format

8-group rota

Grading rubric

Peer-critique form

01 • HOW A SEMINAR SESSION WORKS

Mechanics

Allocation

In Week 1 you rank topic preferences; the instructor allocates one Seminar-1 topic and one Seminar-2 topic per group (published within 3 days), balancing preferences against the presentation rota below. Swaps are allowed only by mutual agreement before Week 2.

The slot: 30' talk + 15' Q&A

The time budget is fixed by the syllabus. Talks are cut at 32 minutes – running over is penalized under "structure & delivery," so rehearse against a timer.

Required talk structure

Every talk – Seminar 1 or 2 – must cover, in this order:

1. **Context & problem** – why this topic matters, what question the papers answer

2. **Survey of main approaches** – the core papers, explained at the mechanism level
3. **Critical analysis** – assumptions, evaluation weaknesses, comparisons across methods (not a summary – this is 25% of your grade)
4. **Demo / example** – a worked example, trace, or live demo when feasible
5. **Relation to projects & open directions** – how it connects to your own project and what remains unsolved

Target 18–22 slides, at least one table comparing methods, and one slide titled “*What we don’t believe / what the papers don’t show.*”

SUBMIT 48 HOURS BEFORE YOUR SLOT

Slides (PDF), a 1–page annotated bibliography (one critical sentence per core + extension paper), a demo link/video if applicable, and an AI-use statement. Everyone else submits reading notes ($\leq \frac{1}{2}$ page) and 2 critical questions 24 hours before, and fills a peer-critique form in class.

Discussant group

Each talk has a designated discussant group (see the rota) who opens the Q&A with two prepared questions. In the rota below, the group presenting the following week discusses the current talk – so read ahead.

02 • WHO PRESENTS WHEN

Presentation rota

Illustrative allocation for a nominal cohort of 8 groups (G1–G8) – the actual rota is published after group registration in Week 1. Groups that present early in Seminar 1 present late in Seminar 2, so preparation time is balanced across the semester.

WEEK	TALK A	TALK B	DISCUSSANTS
W3	G1 – S1-1 Reasoning & Planning	G2 – S1-2 Tool Use & MCP	G3, G4

WEEK	TALK A	TALK B	DISCUSSANTS
W4	G3 – S1-3 Agent Memory	G4 – S1-4 Retrieval-Augmented Agents	G5, G6
W5	G5 – S1-5 Coding Agents	G6 – S1-6 Self-improvement & Reflection	G7, G8
W6	G7 – S1-7 Design Patterns & Orchestration	G8 – S1-8 Evaluation & Benchmarking	G1, G2
W8	G5 – S2-1 Multi-Agent Systems	G6 – S2-2 Multi-Agent Communication	G7, G8
W9	G7 – S2-3 Prompt Injection & Attacks	G8 – S2-4 Advanced Security & Isolation	G1, G2
W10	G1 – S2-5 Safety & Alignment	G2 – S2-6 RL for Agents	G3, G4
W11	G3 – S2-7 Computer-use / Embodied Agents	G4 – S2-8 Production, Governance & Ethics	G5, G6

03 • SEMINAR 1 – WEEKS 3-6

Core topics

S1-1 Advanced Reasoning & Planning

ReAct, Reflexion, Tree/Graph-of-Thoughts, LATS, and hierarchical planning – the paradigms that turn a single LLM call into a reasoning process.

CORE READINGS

- Wei et al. – [Chain-of-Thought Prompting](#) (NeurIPS 2022)
- Yao et al. – [ReAct: Synergizing Reasoning and Acting](#) (ICLR 2023)
- Shinn et al. – [Reflexion](#) (NeurIPS 2023)
- Yao et al. – [Tree of Thoughts](#) (NeurIPS 2023)

EXTENSION (PRESENTING GROUP)

- Besta et al. – [Graph of Thoughts](#) (AAAI 2024)
- Zhou et al. – [Language Agent Tree Search / LATS](#) (ICML 2024)
- Kambhampati et al. – [LLM-Modulo planning frameworks](#) (ICML 2024)

MUST EXPLAIN IN YOUR TALK

- Why intermediate CoT tokens actually help – and why it's an emergent, scale-dependent capability, not a universal fix
- ReAct's Thought→Action→Observation loop, and what it assumes about action cost and reversibility
- How Reflexion differs mechanically from ReAct: episodic memory of reflections, and what feedback signal it needs
- ToT/GoT: what the search actually searches over, and where the evaluator becomes the bottleneck
- The planning-vs-decomposition distinction – why "the model produced steps" ≠ "the model can plan"
- A cost/accuracy comparison across at least 3 of these techniques (tokens, latency, or \$ per task)

Be ready to defend: *What does ReAct assume about observation cost and reversibility?*
Where does Tree-of-Thoughts' gain actually come from – search, or the evaluator?

S1-2 Tool Use & MCP

Function calling and the Model Context Protocol at the protocol level – not "how to build an MCP server," but what MCP standardizes and where its security surface is.

CORE READINGS

- Schick et al. – [Toolformer](#) (NeurIPS 2023)
- Qin et al. – [ToolLLM](#) (ICLR 2024)
- Anthropic – [Introducing the Model Context Protocol](#) (2024 launch) + [current MCP architecture and spec](#)

EXTENSION (PRESENTING GROUP)

- Patil et al. – [Gorilla](#) (NeurIPS 2024)
- [Berkeley Function-Calling Leaderboard](#)
- Google – [Agent2Agent Protocol / A2A](#) (2025)

MUST EXPLAIN IN YOUR TALK

- The mechanics of native function calling: schema → constrained generation → execution → observation

- Why tool selection at scale (thousands of APIs) is a distinct problem from calling one tool correctly
- MCP's host/client/server architecture; tools, resources, and prompts; Elicitation; and current capability discovery
- What changed since the 2024 launch protocol: stateless data layer, [server/discover](#), and deprecated Sampling
- What MCP standardizes that ad-hoc function calling did not, and where the $M \times N \rightarrow M + N$ simplification still leaves integration work
- Why a tool's description and result are untrusted input from the model's perspective
- A2A vs. MCP: agent-to-agent vs. agent-to-tool, and why they're complementary, not competing

Be ready to defend: *Which MCP input creates the largest injection surface, and why?*
What does MCP standardize that plain function calling did not?

S1-3 Agent Memory

Short/long-term and episodic/semantic memory architectures – MemGPT's paging model, Generative Agents' retrieval scoring, and what "forgetting" should mean.

CORE READINGS

- Packer et al. – [MemGPT](#) (2023)
- Park et al. – [Generative Agents](#) (UIST 2023)

EXTENSION (PRESENTING GROUP)

- Sumers et al. – [CoALA memory taxonomy](#) (TMLR 2024)
- Zhong et al. – [MemoryBank](#) (AAAI 2024)
- Liu et al. – [Lost in the Middle](#) (TACL 2024)
- Xu et al. – [A-MEM](#) (2025)

MUST EXPLAIN IN YOUR TALK

- The distinction between context-window ("short-term") and persistent ("long-term") memory
- MemGPT's paging/virtual-context mechanism: what triggers eviction, and what the write policy is
- Generative Agents' retrieval scoring (recency/importance/relevance) and how it was evaluated
- CoALA's memory taxonomy (working vs. long-term; episodic/semantic/procedural)
- Why "lost in the middle" means position matters independently of context capacity

- How memory can be poisoned, and what a write-time or read-time defense would look like

Be ready to defend: *What is the write policy in MemGPT – who decides what is remembered?*

How can memory be poisoned, and what would you log to detect it?

S1-4 Retrieval-Augmented & Knowledge-grounded Agents

RAG and GraphRAG as tools inside an agent loop – and when retrieval actively hurts rather than helps.

CORE READINGS

- Lewis et al. – [Retrieval-Augmented Generation](#) (NeurIPS 2020)
- Edge et al. – [GraphRAG](#) (2024)

EXTENSION (PRESENTING GROUP)

- Asai et al. – [Self-RAG](#) (ICLR 2024)
- Jiang et al. – [Active Retrieval / FLARE](#) (EMNLP 2023)
- Gao et al. – [RAG survey](#) (2024)

MUST EXPLAIN IN YOUR TALK

- The RAG mechanism: retrieve, then condition generation on the retrieved passages
- Naive RAG's failure on "global" questions, and how GraphRAG's community-summary approach addresses it
- Self-RAG: what the reflection tokens actually decide (whether to retrieve, and how to critique retrieved content)
- FLARE: retrieval triggered by the model's own low-confidence tokens, not a fixed schedule
- At least one concrete case where retrieval measurably hurts accuracy
- How retrieval fits as one tool inside an agent's decision loop, not a separate pipeline stage

Be ready to defend: *When does retrieval reduce accuracy?*

How should an agent decide whether to retrieve at all?

S1-5 Coding Agents & Automated SE

SWE-bench and SWE-agent as research objects: agent-computer interface design, evaluation validity, and contamination – not tutorial use of an AI IDE.

CORE READINGS

- Jimenez et al. – [SWE-bench](#) (ICLR 2024)

- Yang et al. – [SWE-agent](#) (NeurIPS 2024)

EXTENSION (PRESENTING GROUP)

- Wang et al. – [OpenHands](#) (ICLR 2025)
- Xia et al. – [Agentless](#) (2024)
- Xia & Zhang – [ChatRepair](#) (ISSTA 2024)

MUST EXPLAIN IN YOUR TALK

- SWE-bench's task format (real GitHub issues + hidden tests) and why "Verified" exists (contamination/flakiness filtering)
- The agent-computer interface in SWE-agent: what custom tools it gives the model, and why that matters for the result
- Agentless's three-phase (localize/repair/validate) pipeline as a counter-argument to "agents need autonomy"
- What SWE-bench does *not* measure (maintainability, security, code-review quality)
- ChatRepair's conversational repair loop: how test-failure feedback shapes the next patch attempt
- One concrete ablation or contamination concern for any of these systems

Be ready to defend: *What does SWE-bench measure and not measure? Why does "Verified" exist?*

Is the agent-computer interface or the underlying model responsible for SWE-agent's gain?

S1-6 Self-improvement & Reflection

Reflexion, Self-Refine, Self-Consistency – and the uncomfortable result that self-correction without an external verifier is weak.

CORE READINGS

- Madaan et al. – [Self-Refine](#) (NeurIPS 2023)
- Wang et al. – [Self-Consistency](#) (ICLR 2023)

EXTENSION (PRESENTING GROUP)

- Huang et al. – [LLMs Cannot Self-Correct Reasoning Yet](#) (ICLR 2024)
- Lightman et al. – [Let's Verify Step by Step](#) (ICLR 2024)
- Zelikman et al. – [STaR](#) (NeurIPS 2022)

MUST EXPLAIN IN YOUR TALK

- Self-Refine's loop: the same model generates, critiques, and revises – no external signal

- Self-Consistency: sampling + majority vote, and why that's a different mechanism from refinement
- Huang et al.'s finding: why self-correction without external feedback often fails or hurts accuracy
- The outcome-reward vs. process-reward distinction, and why process supervision helped in "Let's Verify Step by Step"
- STaR's bootstrap loop: generate rationales, keep the ones reaching correct answers, fine-tune
- Where each technique's feedback signal actually comes from (self vs. external) – the key critical question for this topic

Be ready to defend: *Under what conditions does self-refinement actually work?*

What's the difference between an outcome and a process verifier, and why does it matter for search?

S1-7 Agent Design Patterns & Orchestration

Workflows vs. autonomous agents; prompt chaining, routing, orchestrator-worker, evaluator-optimizer – and why the simplest pattern that works is usually the right one.

CORE READINGS

- Anthropic – [Building Effective Agents](#) (2024)

EXTENSION (PRESENTING GROUP)

- Anthropic – [How we built our multi-agent research system](#) (2025)
- OpenAI – [A Practical Guide to Building Agents](#) (2025)
- [LangGraph documentation](#) – state graphs, checkpointing, human-in-the-loop

MUST EXPLAIN IN YOUR TALK

- The workflow-vs-agent distinction: fixed code path vs. model-directed control flow
- At least three concrete patterns (prompt chaining, routing, orchestrator-worker, evaluator-optimizer) and what each buys/costs
- Why Anthropic argues for the simplest pattern that works, and what triggers moving to a more complex one
- What "durable execution" / checkpointing means for a long-running agent's failure recovery
- How the multi-agent research system's lead/subagent split changes cost and latency, not just capability
- A design-pattern decision your own project actually made, and the alternative you rejected

Be ready to defend: *Why does Anthropic recommend against agents for many tasks?*

S1-8 Agent Evaluation & Benchmarking

AgentBench, GAIA, WebArena, τ -bench, and the recurring pitfalls: contamination, judge bias, and cost-blind leaderboards. Directly useful for your project's evaluation.

CORE READINGS

- Liu et al. – [AgentBench](#) (ICLR 2024)
- Mialon et al. – [GAIA](#) (ICLR 2024)
- Zhou et al. – [WebArena](#) (ICLR 2024)

EXTENSION (PRESENTING GROUP)

- Yao et al. – [\$\tau\$ -bench](#) (2024)
- Kapoor et al. – [AI Agents That Matter](#) (2024)
- Zheng et al. – [Judging LLM-as-a-Judge](#) (NeurIPS 2023)

MUST EXPLAIN IN YOUR TALK

- What AgentBench/GAIA/WebArena each actually measure, and on what kind of task
- $\text{pass}@k$ vs. pass^k : what each measures, and why pass^k is the stricter, deployment-relevant one
- τ -bench's reliability finding: why even strong agents fail pass^8 despite a decent $\text{pass}@1$
- The "AI Agents That Matter" critique: cost-blind leaderboards, inadequate holdouts, conflated benchmarking needs
- LLM-as-judge's known biases (verbosity, position, self-enhancement) and how to audit for them
- Contamination as a threat to any benchmark number, and how you'd test for it

Be ready to defend: *Why report pass^k instead of $\text{pass}@k$?*

What are the failure modes of LLM-as-judge, and how would you audit one?

S2-1 Multi-Agent Systems & Orchestration

AutoGen, MetaGPT, CAMEL – and when a team of agents genuinely beats one agent with better tools.

CORE READINGS

- Wu et al. – [AutoGen](#) (COLM 2024)
- Hong et al. – [MetaGPT](#) (ICLR 2024)
- Li et al. – [CAMEL](#) (NeurIPS 2023)

EXTENSION (PRESENTING GROUP)

- Qian et al. – [ChatDev](#) (ACL 2024)
- Cemri et al. – [Why Do Multi-Agent LLM Systems Fail?](#) (2025)

MUST EXPLAIN IN YOUR TALK

- What AutoGen/MetaGPT/CAMEL actually add structurally (roles, conversation patterns) vs. a single agent with tools
- At least 2–3 concrete failure categories from Cemri et al.'s taxonomy (system design, inter-agent misalignment, task verification)
- A fair, cost-matched comparison design: multi-agent vs. single-agent on the same task and budget
- When a single agent with better tools plausibly wins over a multi-agent team
- ChatDev's chat-chain structure: how communication phases map onto SE lifecycle stages

Be ready to defend: *Is multi-agent buying capability, context isolation, or parallelism?*

How do you evaluate a multi-agent system fairly against a single agent at matched cost?

S2-2 Multi-Agent Communication & Collaboration

Debate, role-play, and negotiation protocols between agents – and where emergent behaviour is real vs. an artifact of prompting.

CORE READINGS

- Park et al. – [Generative Agents](#) (UIST 2023)
- Du et al. – [Multiagent Debate](#) (ICML 2024)

EXTENSION (PRESENTING GROUP)

- Google – [A2A protocol](#) (2025)
- Liang et al. – [Divergent thinking via multi-agent debate](#) (EMNLP 2024)

MUST EXPLAIN IN YOUR TALK

- Multiagent Debate's mechanism: multiple instances propose and critique each other's answers over rounds
- Whether debate gains hold up against self-consistency at matched token budget – the core critical question
- Generative Agents' emergent behavior: what's genuinely emergent vs. an artifact of the simulation's prompting
- What A2A standardizes for agent-to-agent messaging that MCP does not
- How agents could collude, and what a detection/monitoring signal would look like

Be ready to defend: *At equal token budget, does debate actually beat self-consistency?*
How could agents collude, and how would you detect it?

S2-3 Prompt Injection & Attacks

Direct and indirect prompt injection as a *system* vulnerability, not a model one – the OWASP LLM Top 10 and why prompt-level detectors fail.

CORE READINGS

- Greshake et al. – [Indirect Prompt Injection](#) (AISec 2023)
- Perez & Ribeiro – [Ignore Previous Prompt](#) (2022)
- [OWASP – Top 10 for LLM Applications](#) (2025)

EXTENSION (PRESENTING GROUP)

- Zou et al. – [Universal adversarial attacks / GCG](#) (2023)
- Liu et al. – [Formalizing prompt injection](#) (USENIX Security 2024)
- Willison – [The lethal trifecta for AI agents](#) (2025)

MUST EXPLAIN IN YOUR TALK

- The direct/indirect injection distinction, and why indirect is the harder, more dangerous case
- A concrete indirect-injection scenario for your **own project's** tools – not a generic example
- Why "ignore previous instructions"-style detectors are fundamentally unreliable
- GCG's mechanism at a conceptual level: gradient-guided suffix search, and why it transfers across models
- The OWASP LLM Top 10 categories relevant to agents specifically, not just chatbots
- Willison's "lethal trifecta": which of your own project's tools puts it at risk

Test injection and jailbreak techniques only against your own project's system or sanctioned benchmarks (e.g., AgentDojo). Never test against third-party services.

Be ready to defend: *Which of your own project's tools could become an exfiltration channel?*

S2-4 Advanced Security & Isolation

Memory poisoning, multi-agent collusion, tool-supply-chain attacks, and defenses by design – sandboxing, least privilege, and information-flow control.

CORE READINGS

- Chen et al. – [AgentPoison](#) (NeurIPS 2024)
- DeBenedetti et al. – [AgentDojo](#) (NeurIPS 2024, D&B)

EXTENSION (PRESENTING GROUP)

- DeBenedetti et al. – [Defeating Prompt Injections by Design / CaMeL](#) (2025)
- Beurer-Kellner et al. – [Design Patterns for Securing LLM Agents](#) (2025)

MUST EXPLAIN IN YOUR TALK

- AgentPoison's mechanism: poisoning memory/RAG stores so a trigger retrieves a malicious demonstration
- AgentDojo's evaluation design: utility vs. attack-success-rate as two separate axes, not one
- CaMeL's core idea: separating control flow from data flow so untrusted data can't redirect execution
- The utility cost of CaMeL's defense, with concrete numbers – security is not free
- At least 2 of Beurer-Kellner et al.'s design patterns and what each one rules out
- The minimum permission model you'd apply to your own project's tools

Be ready to defend: *What does AgentDojo measure, and what is the utility cost of the best defenses?*

What is the minimum permission model for your project's tools?

S2-5 Safety & Alignment

Guardrails, Constitutional AI, human-in/on-the-loop, and action-boundary enforcement – model alignment vs. system safety.

CORE READINGS

- Bai et al. – [Constitutional AI](#) (2022)
- Ouyang et al. – [RLHF](#) (NeurIPS 2022)

EXTENSION (PRESENTING GROUP)

- Shavit et al. – [Practices for Governing Agentic AI Systems](#) (2023)
- Greenblatt et al. – [Alignment Faking in LLMs](#) (2024)
- Meinke et al. – [In-context Scheming](#) (2024)

MUST EXPLAIN IN YOUR TALK

- Constitutional AI's mechanism: self-critique against a written constitution, no human harm labels
- The model-alignment vs. system-safety distinction, with one concrete example of each failing independently
- Alignment faking: what behavior was observed, and why it complicates "the model says it's aligned" as evidence
- In-context scheming: at least one concrete strategic behavior models exhibited (disabling oversight, exfiltration)
- What an "action boundary" or escalation policy would need to look like for your own project

Be ready to defend: *Which failures does model alignment address, and which need system-level safety?*

What is an "action boundary," and how do you enforce it outside the prompt?

S2-6 Reinforcement Learning for Agents

Verifiable rewards, GRPO, and reasoning-RL (o1/R1-style) – plus reward hacking as the recurring failure mode.

CORE READINGS

- Sutton & Barto – [Reinforcement Learning: An Introduction](#) (ch. 1, 3, 13)
- Ouyang et al. – [RLHF](#) (NeurIPS 2022)
- Shao et al. – [GRPO / DeepSeekMath](#) (2024)

EXTENSION (PRESENTING GROUP)

- DeepSeek-AI – [DeepSeek-R1](#) (2025)
- Lightman et al. – [Process reward models](#) (ICLR 2024)
- Pan et al. – [SWE-Gym](#) (2024)

MUST EXPLAIN IN YOUR TALK

- Why verifiable rewards (not human preference alone) unlocked reasoning-RL at scale
- GRPO's mechanism relative to PPO: what it removes (the critic) and why that matters for cost
- DeepSeek-R1's training recipe at a high level: what changes between R1-Zero and R1
- A concrete reward-hacking example for a tool-using agent – not a generic "reward hacking exists" claim
- SWE-Gym's role: why a training environment with executable verification, not just more data, was the bottleneck

Be ready to defend: *Why did verifiable rewards unlock reasoning-RL?*

Give a reward-hacking example for a tool-using agent.

S2-7 Computer-use / GUI / Embodied Agents

Web agents, computer-use agents, and embodied agents – where perception-action grounding still fails on long horizons.

CORE READINGS

- Wang et al. – [Voyager](#) (TMLR 2024)
- Zhou et al. – [WebArena](#) (ICLR 2024)

EXTENSION (PRESENTING GROUP)

- Xie et al. – [OSWorld](#) (NeurIPS 2024)
- Deng et al. – [Mind2Web](#) (NeurIPS 2023)
- Ahn et al. – [SayCan](#) (CoRL 2022)
- Brohan et al. – [RT-2](#) (CoRL 2023)
- Anthropic – [Developing a computer use model](#) (2024)

MUST EXPLAIN IN YOUR TALK

- Voyager's three components (curriculum, skill library, iterative prompting) and what each solves
- Why OSWorld/WebArena success rates remain low – name a specific failure mode, not "it's hard"
- The perception-grounding problem: screenshots/pixels vs. accessibility trees/DOM as different action interfaces
- What safety control an irreversible real-world action would need before an agent is allowed to take it

Be ready to defend: *Why are OSWorld success rates still low – what actually fails?*

S2-8 Production, Governance & Ethics

Observability, reliability, cost, and governance for non-deterministic agent systems – reliability engineering meets agentic AI.

CORE READINGS

- Huyen – [Agents & Building A Generative AI Platform](#) (free, adapted from AI Engineering, 2025)
- Chan et al. – [Visibility into AI Agents](#) (FAccT 2024)

EXTENSION (PRESENTING GROUP)

- Shavit et al. – [Practices for Governing Agentic AI Systems](#) (2023)
- Gabriel et al. – [The Ethics of Advanced AI Assistants](#) (2024)
- [OpenTelemetry – Semantic Conventions for Generative AI](#) (2024–25)

MUST EXPLAIN IN YOUR TALK

- What SLOs/error budgets mean for a system whose outputs are non-deterministic
- The three visibility measures from Chan et al. (agent identifiers, real-time monitoring, activity logging) and which your project implements
- What OpenTelemetry's GenAI semantic conventions standardize (spans for model calls / agent calls)
- Cost-per-successful-task as a metric, and why raw accuracy alone is insufficient for a governance conversation
- One concrete governance or ethical risk specific to your own project's deployment context

Be ready to defend: *Which SLOs make sense for a non-deterministic agent?*

What must be logged for an audit trail (agent IDs, tool calls, approvals)?

CRITERION	WEIGHT	WHAT EARNS A HIGH SCORE
Scientific accuracy & depth	30%	Core + extension papers correct in detail; mechanisms explained at algorithm level; recent results included
Critical analysis	30%	Identifies assumptions, evaluation weaknesses, and contradictions between papers; a supported comparison table; your own position, argued
Structure & delivery (30')	20%	Required structure followed; clear narrative; on time; slides readable; demo works
Q&A & defense (15')	20%	Answers precisely, concedes limits, argues with evidence; every member can answer

INDIVIDUAL GRADING

All four criteria are scored per group, except that Q&A performance can move an **individual's** score within the group: if you cannot defend your own slides, expect a lower personal mark even if your group's talk was strong. There is no separate graded "critique of other groups" line – see the participation bonus below for how asking good questions is rewarded instead.

PARTICIPATION BONUS – ASKING QUESTIONS, BOTH ROUNDS

Asking a substantive, evidence-based question during Q&A – as an assigned **discussant** or simply from the floor during any other group's talk – earns participation credit, tracked from live Q&A plus your peer-critique forms. Accumulated across Seminar 1 and Seminar 2, this converts into a **bonus of up to +0.5 points added to your final course grade** (10-point scale) – on top of the 25/25/50 weighting, not a substitute for it, and capped so your final grade can never exceed 10/10.

As a non-presenting student, submit one peer-critique form per talk:

- One claim in the talk you found **well-supported**, and why (2–3 sentences)
- One claim you found **unsupported or wrong**, with the evidence or paper that contradicts it
- One question you would ask the **authors of the core paper** – not the presenters
- Scores 1–5 on: accuracy · critical depth · delivery · usefulness for your own project

As the **discussant group** for a talk, prepare two questions in advance and open the Q&A with them. This form is also what the participation bonus above is based on – it is no longer a separately graded rubric line, but asking good questions still pays off.

← BACK

Course overview

NEXT →

Group project

C05151 · Advanced Agentic Artificial Intelligence

HCMUT · Faculty of Computer Science and Engineering · Lecturer: Lê Xuân Bách